

```
#import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# read dataset
dataframe = pd.read_csv("Jamboree_Admission.csv")
```

✓ Basic data cleaning and exploration

```
dataframe.head()
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

```
dataframe.describe()
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA
count	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000
mean	250.500000	316.472000	107.192000	3.114000	3.374000	3.48400	8.576440
std	144.481833	11.295148	6.081868	1.143512	0.991004	0.92545	0.604813
min	1.000000	290.000000	92.000000	1.000000	1.000000	1.00000	6.800000
25%	125.750000	308.000000	103.000000	2.000000	2.500000	3.00000	8.127500
50%	250.500000	317.000000	107.000000	3.000000	3.500000	3.50000	8.560000
75%	375.250000	325.000000	112.000000	4.000000	4.000000	4.00000	9.040000

✓ EDA:

Range of Serial No. is 1 to 500. i.e. Unique Row Values.

Range of GRE Score is 290 to 340.

Range of TOEFL Score is 92 to 120.

Range of CGPA is 6.8 to 9.92.

Range of University Rating(Catagorical) is 1 to 5.

Range of SOP and LOR (Catagorical) is 1 to 5.

The dataset has been thoroughly examined, revealing the absence of outliers. All data points fall within the expected range, ensuring the integrity and reliability of the dataset.

```
dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score              500 non-null   int64
2   TOEFL Score            500 non-null   int64
3   University Rating      500 non-null   int64
4   SOP                    500 non-null   float64
5   LOR                    500 non-null   float64
6   CGPA                   500 non-null   float64
7   Research               500 non-null   int64
8   Chance of Admit        500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

```
# Checking NA values
```

```
dataframe.isnull().sum()
```

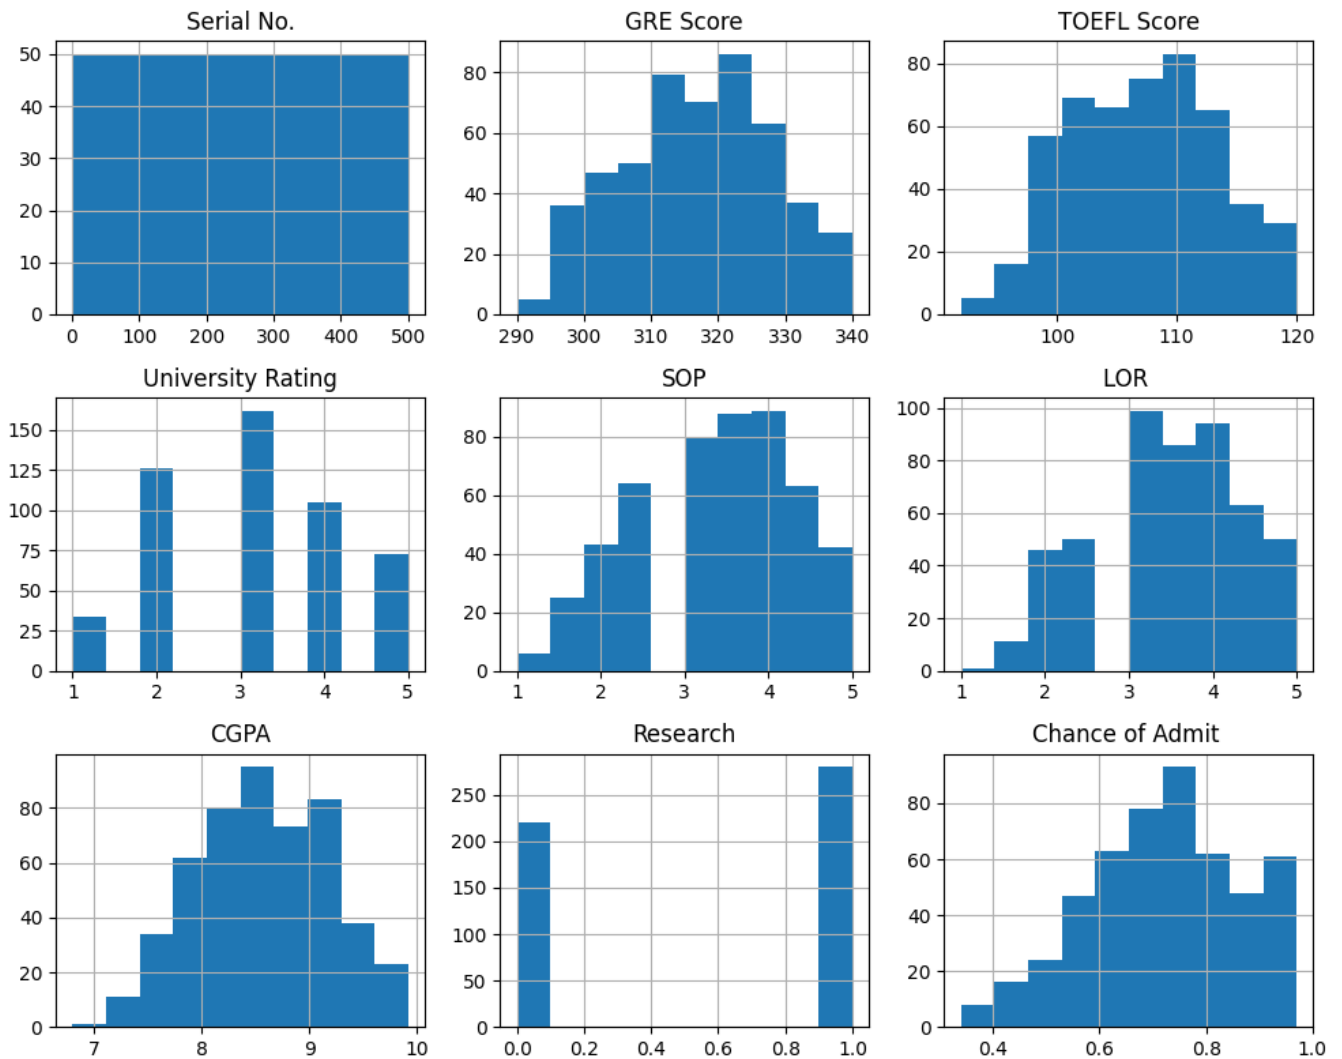
```
Serial No.      0
GRE Score        0
TOEFL Score      0
University Rating 0
SOP              0
LOR              0
CGPA             0
Research         0
Chance of Admit  0
dtype: int64
```

```
# Visualizing the distribution of numerical variables
```

```
dataframe.hist(figsize=(10, 8))
```

```
plt.tight_layout()
```

```
plt.show()
```



```
dataframe.nunique()
```



Serial No.	500
GRE Score	49
TOEFL Score	29
University Rating	5
SOP	9
LOR	9
CGPA	184
Research	2
Chance of Admit	61

dtype: int64


```
# Assuming 'dataframe' is your DataFrame
dataframe.drop(columns=['Serial No.'], inplace=True)
```

```
import seaborn as sns
# Univariate analysis for continuous variables (distribution plots)
continuous_vars = ['GRE Score', 'TOEFL Score', 'CGPA']
for var in continuous_vars:
    plt.figure(figsize=(8, 5))
    sns.histplot(dataframe[var], kde=True)
    plt.title(f'Distribution of {var}')
    plt.xlabel(var)
    plt.ylabel('Frequency')
    plt.show()
```


```
# Univariate analysis for categorical variables (bar plots/count plots)
categorical_vars = ['Research', 'University Rating', 'SOP', 'LOR ', 'Chance of Admit ']
for var in categorical_vars:
    plt.figure(figsize=(8, 5))
    sns.countplot(data=dataframe
                  , x=var)
    plt.title(f'Count Plot of {var}')
    plt.xlabel(var)
    plt.ylabel('Count')
    plt.show()
```



Distribution of CGPA




Count Plot of Research




Count Plot of University Rating








Count Plot of SOP



Count Plot of LOR






✓ EDA

The attributes GRE Scores, TOEFL Scores, University Rating, CGPA, SOP, and LOR exhibit a normal distribution pattern, indicating that the majority of observations cluster around


the mean, with symmetric tails on either side. This distributional characteristic suggests a balanced distribution of values across the dataset, aligning with typical expectations for these attributes.

```
#Bivariate Analysis '  
for var in continuous_vars+categorical_vars[:-1]:  
    if var=='Reaseach':  
        continue  
    plt.figure(figsize=(8, 5))  
    sns.scatterplot(data=dataframe, x=var, y='Chance of Admit ' )  
    plt.title(f'Scatter Plot: {var} vs. Chance of Admit')  
    plt.xlabel(var)  
    plt.ylabel('Chance of Admit')  
    plt.show()
```


Scatter Plot: GRE Score vs. Chance of Admit




Scatter Plot: TOEFL Score vs. Chance of Admit




Scatter Plot: CGPA vs. Chance of Admit







Scatter Plot: Research vs. Chance of Admit



Scatter Plot: University Rating vs. Chance of Admit





- As the GRE score increases, there seems to be a positive trend in the 'Chance of Admit'. Higher GRE scores are associated with higher chances of admission.
- Similar to the GRE score, higher TOEFL scores appear to correlate with higher chances of admission. There is a positive relationship between TOEFL scores and the 'Chance of Admit'.
- There is a noticeable positive relationship between CGPA and the 'Chance of Admit'. Higher CGPA scores correspond to higher chances of admission.
- The scatter plot suggests that there may be a positive relationship between the Statement of Purpose (SOP) score and the 'Chance of Admit'. Students with stronger SOP scores tend to have higher chances of admission.
- Similarly, there appears to be a positive correlation between the Letter of Recommendation (LOR) score and the 'Chance of Admit'. Higher LOR scores are associated with higher chances of admission.
- It is seen that students with research experience are more likely to be admitted compared to those without research experience.

✓ Data Preprocessing

```
#Duplicate value check
duplicate_rows = dataframe[dataframe.duplicated()]
dataframe.drop_duplicates(inplace=True)
```




```
#Missing value treatment
dataframe.isnull().sum()
```


```
GRE Score      0
TOEFL Score    0
University Rating  0
SOP            0
LOR            0
CGPA           0
Research       0
Chance of Admit  0
dtype: int64
```

```
# Visualizing distribution of data points using box plots
plt.figure(figsize=(12, 8))
sns.boxplot(data=dataframe[['GRE Score', 'TOEFL Score']])
plt.title('Box Plot of Relevant Columns')
plt.xticks(rotation=45)
plt.show()
```

```
# Visualizing distribution of data points using box plots
plt.figure(figsize=(12, 8))
sns.boxplot(data=dataframe['CGPA'])
plt.title('Box Plot of Relevant Columns')
plt.xticks(rotation=45)
plt.show()
```

```
# Visualizing distribution of data points using box plots
plt.figure(figsize=(12, 8))
sns.boxplot(data=dataframe[categorical_vars[:-1]])
plt.title('Box Plot of Relevant Columns')
plt.xticks(rotation=45)
plt.show()
```





```
#Outlier treatment
for var in continuous_vars+categorical_vars:
    Q1 = dataframe[var].quantile(0.25)
    Q3 = dataframe[var].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    dataframe[var] = dataframe[var].clip(lower=lower_bound, upper=upper_bound)
```

```
# Feature Engineering
dataframe['strength'] = dataframe['SOP'] + dataframe['LOR ']
dataframe.drop(columns=['SOP', 'LOR '], inplace=True)
dataframe['strength']
```

```
0      9.0
1      8.5
2      6.5
3      6.0
4      5.0
...
495     8.5
496    10.0
497     9.5
498     9.0
499     9.0
Name: strength, Length: 500, dtype: float64
```

```
dataframe
```

	GRE Score	TOEFL Score	University Rating	CGPA	Research	Chance of Admit	strength
0	337	118	4	9.65	1	0.92	9.0
1	324	107	4	8.87	1	0.76	8.5
2	316	104	3	8.00	1	0.72	6.5
3	322	110	3	8.67	1	0.80	6.0
4	314	103	2	8.21	0	0.65	5.0
...
495	332	108	5	9.02	1	0.87	8.5
496	337	117	5	9.87	1	0.96	10.0
497	330	120	5	9.56	1	0.93	9.5
498	312	103	4	8.43	0	0.73	9.0
499	327	113	4	9.04	0	0.84	9.0

500 rows × 7 columns

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder

# Split data into features (X) and target variable (y)
X = dataframe.drop(columns=['Chance of Admit '])
y = dataframe['Chance of Admit ']
columns = X.columns
# Normalize numerical features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=

```

X_train

```

array([[ 0.40128156,  0.62675052, -0.09979274,  0.41965703,  0.88640526,
         0.36760129],
       [-0.0418297 ,  0.62675052,  0.77558214, -0.06031039, -1.12815215,
         1.22782272],
       [-1.193919 , -0.8545401 , -0.09979274, -0.12651279, -1.12815215,
        -0.20587966],
       ...,
       [-1.28254125, -1.34830364, -1.85054249, -2.19533785, -1.12815215,
        -1.63958204],
       [-0.66218548, -0.36077656, -0.97516761, -1.48366203, -1.12815215,
        -0.20587966],

```



```
[-0.21907421, -0.19618871, -0.97516761, -0.5402778 , -1.12815215,
 0.36760129]])
```

✓ Model building

```
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
# Build Linear Regression model
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)

# Predictions train data
y_train_pred = linear_reg.predict(X_train)

coefficients = linear_reg.coef_
coefficients

array([0.02724941, 0.0171936 , 0.00170568, 0.06804298, 0.01188039,
       0.01769839])

# Display model coefficients with column names
coefficients = linear_reg.coef_
coefficients_df = pd.DataFrame({'Feature': columns, 'Coefficient': coefficients})
print("\nLinear Regression Coefficients:")
print(coefficients_df)

Linear Regression Coefficients:
   Feature  Coefficient
0  GRE Score    0.027249
1  TOEFL Score  0.017194
2  University Rating 0.001706
3         CGPA    0.068043
4     Research    0.011880
5    strength    0.017698

# Model statistics
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
train_r2 = r2_score(y_train, y_train_pred)
train_rmse, train_r2

(0.05960406654108408, 0.8195686865143701)
```

✓ Root mean square error and r2 are good for train dataset.

```
# Lasso Regression
lasso_reg = Lasso(alpha=1) # The alpha parameter for regularization strength
lasso_reg.fit(X_train, y_train)
lasso_train_rmse = np.sqrt(mean_squared_error(y_train, lasso_reg.predict(X_train)))
lasso_train_r2 = r2_score(y_train, lasso_reg.predict(X_train))

print("Lasso Regression Model Statistics:")
print("Train RMSE:", lasso_train_rmse)
print("Train R^2 Score:", lasso_train_r2)

Lasso Regression Model Statistics:
Train RMSE: 0.1403201161630078
Train R^2 Score: 0.0

# Ridge Regression
ridge_reg = Ridge(alpha=0.5) # You can adjust the alpha parameter for regularization strength
ridge_reg.fit(X_train, y_train)
ridge_train_rmse = np.sqrt(mean_squared_error(y_train, ridge_reg.predict(X_train)))
ridge_train_r2 = r2_score(y_train, ridge_reg.predict(X_train))

print("Ridge Regression Model Statistics:")
print("Train RMSE:", ridge_train_rmse)
print("Train R^2 Score:", ridge_train_r2)

Ridge Regression Model Statistics:
Train RMSE: 0.05960422545786763
Train R^2 Score: 0.8195677243786006
```

✓ Testing the assumptions of the linear regression model

```
#Multicollinearity check by VIF score
from statsmodels.stats.outliers_influence import variance_inflation_factor
def calculate_vif(X):
    vif_data = pd.DataFrame()
    vif_data["Feature"] = columns
    vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]
    return vif_data

# Convert training DataFrame if they are numpy arrays
X_train_df = pd.DataFrame(X_train)
vif_scores_train = calculate_vif(X_train_df)
print("VIF Scores for Training Data:")
print(vif_scores_train)

VIF Scores for Training Data:
      Feature  VIF
0  GRE Score  4.488243
1  TOEFL Score  3.637750
```

2	University Rating	2.528933
3	CGPA	4.653647
4	Research	1.517161
5	strength	2.719518

- There is some degree of multicollinearity present in the dataset, with several variables exhibiting moderate levels of multicollinearity, none of the VIF scores exceed the commonly accepted threshold of 5, suggesting that multicollinearity may not be a significant issue in this model.

```
# Calculate residuals
residuals = y_train - y_train_pred


# Calculate mean of residuals
mean_residuals = np.mean(residuals)

print("Mean of Residuals:", mean_residuals)
```

Mean of Residuals: -5.620504062164855e-17

- The mean of residuals is close to zero, indicating that the model's predictions are not having multicollinearity and unbiased on average.

```
# Create residual plot
plt.scatter(y_train_pred, residuals)
plt.axhline(y=0, color='r')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()
```



- ✓ The absence of any discernible pattern in the residual plot suggests that multicollinearity among the predictor variables is not a significant concern.

```
from statsmodels.stats.diagnostic import het_goldfeldquandt
```

```
# Perform Goldfeld-Quandt test
```

```
gq_test = het_goldfeldquandt(residuals, X_train)
```

```
print("Goldfeld-Quandt Test p-value:", gq_test[1])
```

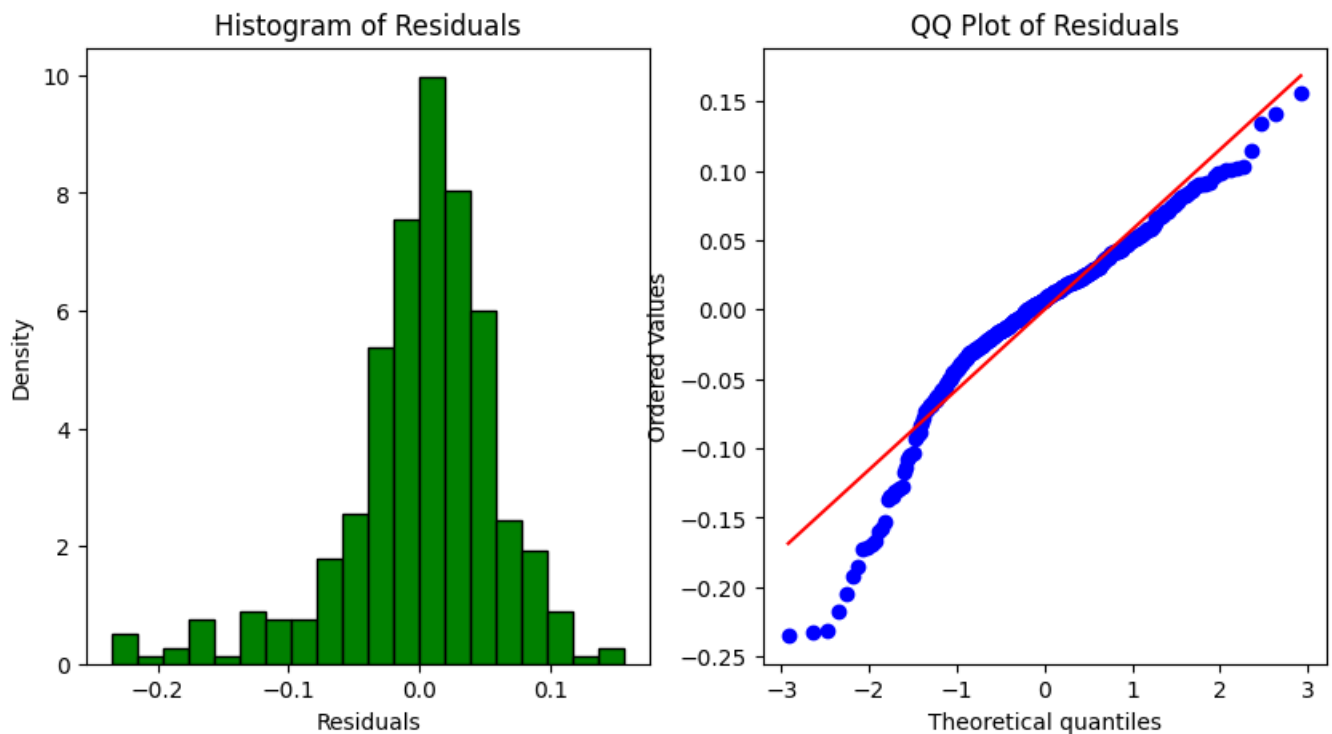
```
Goldfeld-Quandt Test p-value: 0.6848802285877097
```

- ✓ p-value is more than 0.05. So, there is insufficient evidence to conclude that heteroscedasticity is present in the residuals

```
import scipy.stats as stats
# Plot histogram/density plot of residuals
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.hist(residuals, bins=20, density=True, color='g', edgecolor='black')
plt.title('Histogram of Residuals')
plt.xlabel('Residuals')
plt.ylabel('Density')

# Plot QQ plot of residuals
plt.subplot(1, 2, 2)
stats.probplot(residuals, dist="norm", plot=plt)
plt.title('QQ Plot of Residuals')
plt.show()

# Assess normality
normality_result = stats.normaltest(residuals)
print("Normality Test p-value:", normality_result.pvalue)
```



Normality Test p-value: 3.9414129522879716e-20

✓ Residuals may visually appear to be in almost a normal distribution.

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_train, y_train_pred)

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))

# Calculate R-squared (R2)
r2 = r2_score(y_train, y_train_pred)

# Calculate Adjusted R-squared (Adj R2)
n = len(y_train)
p = X.shape[1]
adj_r2 = 1 - ((1 - r2) * (n - 1) / (n - p - 1))

print("Mean Absolute Error (MAE):", mae)
print("Root Mean Squared Error (RMSE):", rmse)
print("R-squared (R2):", r2)
print("Adjusted R-squared (Adj R2):", adj_r2)
```

```
Mean Absolute Error (MAE): 0.042542909235852445
Root Mean Squared Error (RMSE): 0.05960406654108408
R-squared (R2): 0.8195686865143701
Adjusted R-squared (Adj R2): 0.8168140099726048
```

The MAE and RMSE values are relatively low, suggesting that the model's predictions are close to the actual values. The R-squared and Adjusted R-squared values are relatively high, indicating that a large proportion of the variance in the target variable is explained by the model.

```
# Train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on test data
y_test_pred = model.predict(X_test)
```

```
# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_test_pred)

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))
```